

# Explainable AutoML - Titanic Survival Classification Demo

```
In [1]: # Author Hussain Abbas
# Copyright © Stats AI 2021. All Rights Reserved

import tensorflow as tf
import autokeras as ak
from tensorflow.keras import backend as K
import keras_tuner
import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
import re
from sklearn.model_selection import cross_val_score, KFold, train_test_split
from sklearn.metrics import roc_auc_score, precision_score, recall_score, fbeta_score, roc_curve
from sklearn.metrics import confusion_matrix, ConfusionMatrixDisplay
from tqdm import tqdm
```

```
In [2]: # Verify GPU is detected and working
print("Num GPUs Available: ", len(tf.config.list_physical_devices('GPU')))
```

Num GPUs Available: 1

```
In [3]: TRAIN_DATA_URL = "https://storage.googleapis.com/tf-datasets/titanic/train.csv"
TEST_DATA_URL = "https://storage.googleapis.com/tf-datasets/titanic/eval.csv"

#datasets Located in C:/Users/USER/.keras/datasets

train_file_path = tf.keras.utils.get_file("train.csv", TRAIN_DATA_URL)
test_file_path = tf.keras.utils.get_file("eval.csv", TEST_DATA_URL)

train_data = pd.read_csv(train_file_path)
test_data = pd.read_csv(test_file_path)

df = pd.concat([train_data, test_data])
df.drop_duplicates(inplace=True)
df = df.reset_index()
df = df.drop(['index'], axis=1)

print('All Data Summary')
print(df.describe())
print('\n')

print('Train Data Summary')
print(train_data.describe())
print('\n')

print('Test Data Summary')
print(test_data.describe())
```

All Data Summary					
	survived	age	n_siblings_spouses	parch	fare
count	781.000000	781.000000	781.000000	781.000000	781.000000
mean	0.413572	29.622817	0.524968	0.417414	34.750464
std	0.492789	13.764671	0.987592	0.838132	52.237906
min	0.000000	0.420000	0.000000	0.000000	0.000000
25%	0.000000	22.000000	0.000000	0.000000	8.050000
50%	0.000000	28.000000	0.000000	0.000000	15.900000
75%	1.000000	36.000000	1.000000	1.000000	34.020800
max	1.000000	80.000000	8.000000	6.000000	512.329200

Train Data Summary					
	survived	age	n_siblings_spouses	parch	fare
count	627.000000	627.000000	627.000000	627.000000	627.000000
mean	0.387560	29.631308	0.545455	0.379585	34.385399
std	0.487582	12.511818	1.151090	0.792999	54.597730
min	0.000000	0.750000	0.000000	0.000000	0.000000
25%	0.000000	23.000000	0.000000	0.000000	7.895800
50%	0.000000	28.000000	0.000000	0.000000	15.045800
75%	1.000000	35.000000	1.000000	0.000000	31.387500
max	1.000000	80.000000	8.000000	5.000000	512.329200

Test Data Summary					
	survived	age	n_siblings_spouses	parch	fare
count	264.000000	264.000000	264.000000	264.000000	264.000000
mean	0.375000	28.720985	0.469697	0.386364	27.023880
std	0.485042	14.157538	0.978393	0.837775	34.973108
min	0.000000	0.420000	0.000000	0.000000	0.000000
25%	0.000000	21.000000	0.000000	0.000000	7.925000
50%	0.000000	28.000000	0.000000	0.000000	13.250000
75%	1.000000	35.250000	1.000000	0.000000	27.900000
max	1.000000	74.000000	8.000000	6.000000	263.000000

```
In [4]: print('Train Data')
train_data.head()
```

Train Data

Out[4]:

	survived	sex	age	n_siblings_spouses	parch	fare	class	deck	embark_town	alone
0	0	male	22.0	1	0	7.2500	Third	unknown	Southampton	n
1	1	female	38.0	1	0	71.2833	First	C	Cherbourg	n
2	1	female	26.0	0	0	7.9250	Third	unknown	Southampton	y
3	1	female	35.0	1	0	53.1000	First	C	Southampton	n
4	0	male	28.0	0	0	8.4583	Third	unknown	Queenstown	y

```
In [5]: print('Test Data')
test_data.head()
```

Test Data

Out[5]:

	survived	sex	age	n_siblings_spouses	parch	fare	class	deck	embark_town	alone
0	0	male	35.0	0	0	8.0500	Third	unknown	Southampton	y
1	0	male	54.0	0	0	51.8625	First	E	Southampton	y
2	1	female	58.0	0	0	26.5500	First	C	Southampton	y
3	1	female	55.0	0	0	16.0000	Second	unknown	Southampton	y
4	1	male	34.0	0	0	13.0000	Second	D	Southampton	y

```
In [6]: def recall_m(y_true, y_pred):
    true_positives = K.sum(K.round(K.clip(y_true * y_pred, 0, 1)))
    possible_positives = K.sum(K.round(K.clip(y_true, 0, 1)))
    recall = true_positives / (possible_positives + K.epsilon())
    return recall

def precision_m(y_true, y_pred):
    true_positives = K.sum(K.round(K.clip(y_true * y_pred, 0, 1)))
    predicted_positives = K.sum(K.round(K.clip(y_pred, 0, 1)))
    precision = true_positives / (predicted_positives + K.epsilon())
    return precision

def f_beta_score(y_true, y_pred):

    a = 0.5 ** 2
    b = 1 + a

    precision = precision_m(y_true, y_pred)

    recall = recall_m(y_true, y_pred)

    return b*((precision*recall)/(a*precision+recall+K.epsilon()))

def ak_predict(model, data):

    pred_input = data.astype(np.compat.unicode)
    predicted = model.predict(pred_input).flatten()
    pred_result = predicted

    #cut_off = 0.5

    #pred_result = [1 if x > cut_off else 0 for x in predicted]

    return pred_result

def jdl(y_true, y_pred, smooth=100):
    """
    Jaccard = (|X & Y|) / (|X| + |Y| - |X & Y|)
              = sum(|A*B|) / (sum(|A|) + sum(|B|) - sum(|A*B|))

    The jaccard distance loss is usefull for unbalanced datasets. This has been
    shifted so it converges on 0 and is smoothed to avoid exploding or disapearing
    gradient.

    Ref: https://en.wikipedia.org/wiki/Jaccard_index

    @url: https://gist.github.com/wassname/f1452b748efcbeb4cb9b1d059dce6f96
    @author: wassname
    """

    intersection = K.sum(K.abs(y_true * y_pred), axis=-1)
    sum_ = K.sum(K.abs(y_true) + K.abs(y_pred), axis=-1)
    jac = (intersection + smooth) / (sum_ - intersection + smooth)
    return (1 - jac) * smooth
```

[illegible]

```
print("Issue training model")

try:
    # Save the model after each j, i iteration
    model = clf.export_model()
    model.save(model_name, save_format="tf")

except:
    print("Issue saving model")

results = pd.DataFrame(results, columns = ['model_name', 'j', 'i', 'Test_loss', 'Loss:JDL', 'Loss:Binary Cross En
                                         'AUC', 'Accuracy', 'Precision', 'Recall', 'F_Beta_Score'])
```

```
In [13]: #best_model = results.loc[np.argmax(results.test_accuracy)].model_name

best_model = results.loc[np.argmax(results.F_Beta_Score)].model_name

best_model
```

Out[13]: 'model\_autokeras\_0\_0'

```
In [14]: from tensorflow.keras.models import load_model

my_custom_objects={'f_beta_score': f_beta_score,
                   'jdl': jdl, }
my_custom_objects.update(ak.CUSTOM_OBJECTS)

model_ak = load_model(best_model, custom_objects=my_custom_objects)
```

```
In [15]: model_ak.summary()
```

Model: "model"

Layer (type)	Output Shape	Param #
=====		
input_1 (InputLayer)	[(None, 9)]	0
-----		
multi_category_encoding (Mul	(None, 9)	0
-----		
dense (Dense)	(None, 1024)	10240
-----		
batch_normalization (BatchNo	(None, 1024)	4096
-----		
re_lu (ReLU)	(None, 1024)	0
-----		
dropout (Dropout)	(None, 1024)	0
-----		
dropout_1 (Dropout)	(None, 1024)	0
-----		
dense_1 (Dense)	(None, 1)	1025
-----		
classification_head_1 (Activ	(None, 1)	0
=====		
Total params: 15,361		
Trainable params: 13,313		
Non-trainable params: 2,048		
-----		

```
In [16]: # type: pandas.core.frame.DataFrame
pred_input = x_test.astype(np.compat.unicode)

# type: numpy.ndarray
predicted = model_ak.predict(pred_input).flatten()

cut_off = 0.5

pred_result = [1 if x > cut_off else 0 for x in predicted]
pred_result = np.array(pred_result)

actual = y_test.to_numpy()
actual = actual.flatten()

cm = tf.math.confusion_matrix(actual, pred_result)
cm = cm/cm.numpy().sum(axis=1)[:, tf.newaxis]

sns.heatmap(
    cm, annot=True,
    xticklabels=['no', 'yes'],
    yticklabels=['no', 'yes'])
plt.xlabel("Predicted")
plt.ylabel("True")

...

https://towardsdatascience.com/accuracy-precision-recall-or-f1-331fb37c5cb9
https://stackoverflow.com/questions/44172162/f1-score-vs-roc-auc

- Maximize Precision when False Positives are of concern
- Maximize Recall when False Negatives are of concern
- Maximize F1 Score when both are important and classes are unbalanced

...

auc_score = roc_auc_score(actual, pred_result)
precision = precision_score(actual, pred_result)
recall = recall_score(actual, pred_result)
f_beta = fbeta_score(actual, pred_result, beta = 0.5)


print("Cut-Off:", cut_off)
print("ROC-AUC-Score:", auc_score)
print('Precision: ' + str(precision))
print('Recall: ' + str(recall))
print('F_Beta: ' + str(f_beta))


y_test_classes = list(set(y_test))

# print Confusion Matrix from SkLearn
cm = confusion_matrix(actual, pred_result, labels = y_test_classes)

#cm = cm.astype('float') / cm.sum(axis=1)[:, np.newaxis]
disp = ConfusionMatrixDisplay(confusion_matrix=cm, display_labels = y_test_classes)
disp.plot();
```

Cut-Off: 0.5  
ROC-AUC-Score: 0.7501672240802675  
Precision: 0.7735849056603774  
Recall: 0.6307692307692307  
F\_Beta: 0.740072202166065






```
In [17]: # compute true positive rate and false positive rate
false_positive_rate, true_positive_rate, thresholds = roc_curve(actual, pred_result)

# plotting them against each other
def plot_roc_curve(false_positive_rate, true_positive_rate, label=None):
    plt.plot(false_positive_rate, true_positive_rate, linewidth=2, label=label)
    plt.plot([0, 1], [0, 1], 'r', linewidth=4)
    plt.axis([0, 1, 0, 1])
    plt.xlabel('False Positive Rate (FPR)', fontsize=16)
    plt.ylabel('True Positive Rate (TPR)', fontsize=16)

plt.figure(figsize=(14, 7))
plot_roc_curve(false_positive_rate, true_positive_rate)
plt.show()

# https://towardsdatascience.com/predicting-the-survival-of-titanic-passengers-30870ccc7e8
# https://stackoverflow.com/questions/44172162/f1-score-vs-roc-auc
```





```
In [18]: # Naive Random Coin Flip Classifier Performance

predicted = np.random.randint(0,2, size = len(y_test))
pred_result = predicted.flatten()

actual = y_test.to_numpy()
actual = actual.flatten()

cm = tf.math.confusion_matrix(actual, pred_result)
cm = cm/cm.numpy().sum(axis=1)[:, tf.newaxis]

sns.heatmap(
    cm, annot=True,
    xticklabels=['no', 'yes'],
    yticklabels=['no', 'yes'])
plt.xlabel("Predicted")
plt.ylabel("True")

#true_positives = tf.math.count_nonzero(pred_result * actual)
#true_negatives = tf.math.count_nonzero((pred_result - 1) * (actual - 1))
#false_positives = tf.math.count_nonzero(pred_result * (actual - 1))
#false_negatives = tf.math.count_nonzero((pred_result - 1) * actual)

#precision = true_positives / (true_positives + false_positives)
#recall = true_positives / (true_positives + false_negatives)
#f1 = 2 * precision * recall / (precision + recall)

#print("Precision: " + str(np.array(precision).flatten()[0]))
#print("Recall: " + str(np.array(recall).flatten()[0]))
#print("F1: " + str(np.array(f1).flatten()[0]))
#print('')

...

https://towardsdatascience.com/accuracy-precision-recall-or-f1-331fb37c5cb9
https://stackoverflow.com/questions/44172162/f1-score-vs-roc-auc

- Maximize Precision when False Positives are of concern
- Maximize Recall when False Negatives are of concern
- Maximize F1 Score when both are important and classes are unbalanced

...

auc_score = roc_auc_score(actual, pred_result)
precision = precision_score(actual, pred_result)
recall = recall_score(actual, pred_result)
f_beta = fbeta_score(actual, pred_result, beta = 1)

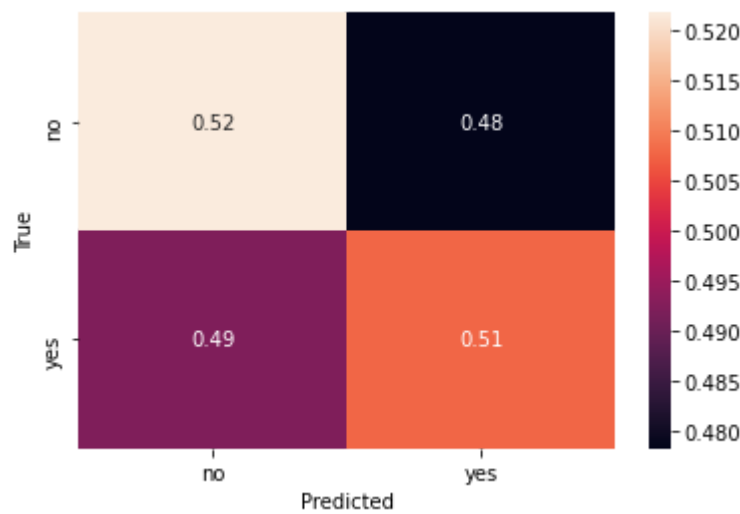
print("ROC-AUC-Score:", auc_score)
print('Precision: ' + str(precision))
print('Recall: ' + str(recall))
print('F_Beta: ' + str(f_beta))

y_test_classes = list(set(y_test))


# print Confusion Matrix from Sklearn
cm = confusion_matrix(actual, pred_result, labels = y_test_classes)

#cm = cm.astype('float') / cm.sum(axis=1)[:, np.newaxis]
disp = ConfusionMatrixDisplay(confusion_matrix=cm, display_labels = y_test_classes)
disp.plot();
```

ROC-AUC-Score: 0.5147157190635451  
Precision: 0.42857142857142855  
Recall: 0.5076923076923077  
F\_Beta: 0.46478873239436613







## Explainable AI using Dalex

```
In [19]: import dalex as dx
```

```
In [20]: X, y = df.drop('survived', axis=1), df.survived
n, p = X.shape
```

```
In [21]: explainer_keras = dx.Explainer(model_ak,
                                         data = X,
                                         y = y,
                                         predict_function = ak_predict,
                                         label = 'autokeras',
                                         #predict_function = dx._explainer.yhat.yhat_tf_classification,
                                         model_type = 'classification'
                                         )
```

Preparation of a new explainer is initiated

```
-> data          : 781 rows 9 cols
-> target variable : Parameter 'y' was a pandas.Series. Converted to a numpy.ndarray.
-> target variable : 781 values
-> model_class    : tensorflow.python.keras.engine.functional.Functional (default)
-> label          : autokeras
-> predict function : <function ak_predict at 0x000001D1EDEE6F70> will be used
-> predict function : Accepts pandas.DataFrame and numpy.ndarray.
-> predicted values : min = 0.134, mean = 0.458, max = 1.0
-> model type     : classification will be used
-> residual function : difference between y and yhat (default)
-> residuals      : min = -0.782, mean = -0.0446, max = 0.758
-> model_info     : package tensorflow
```

A new explainer has been created!

```
In [22]: explainer_keras.model_performance()
```

Out[22]:

	recall	precision	f1	accuracy	auc
autokeras	0.594427	0.803347	0.683274	0.772087	0.827971


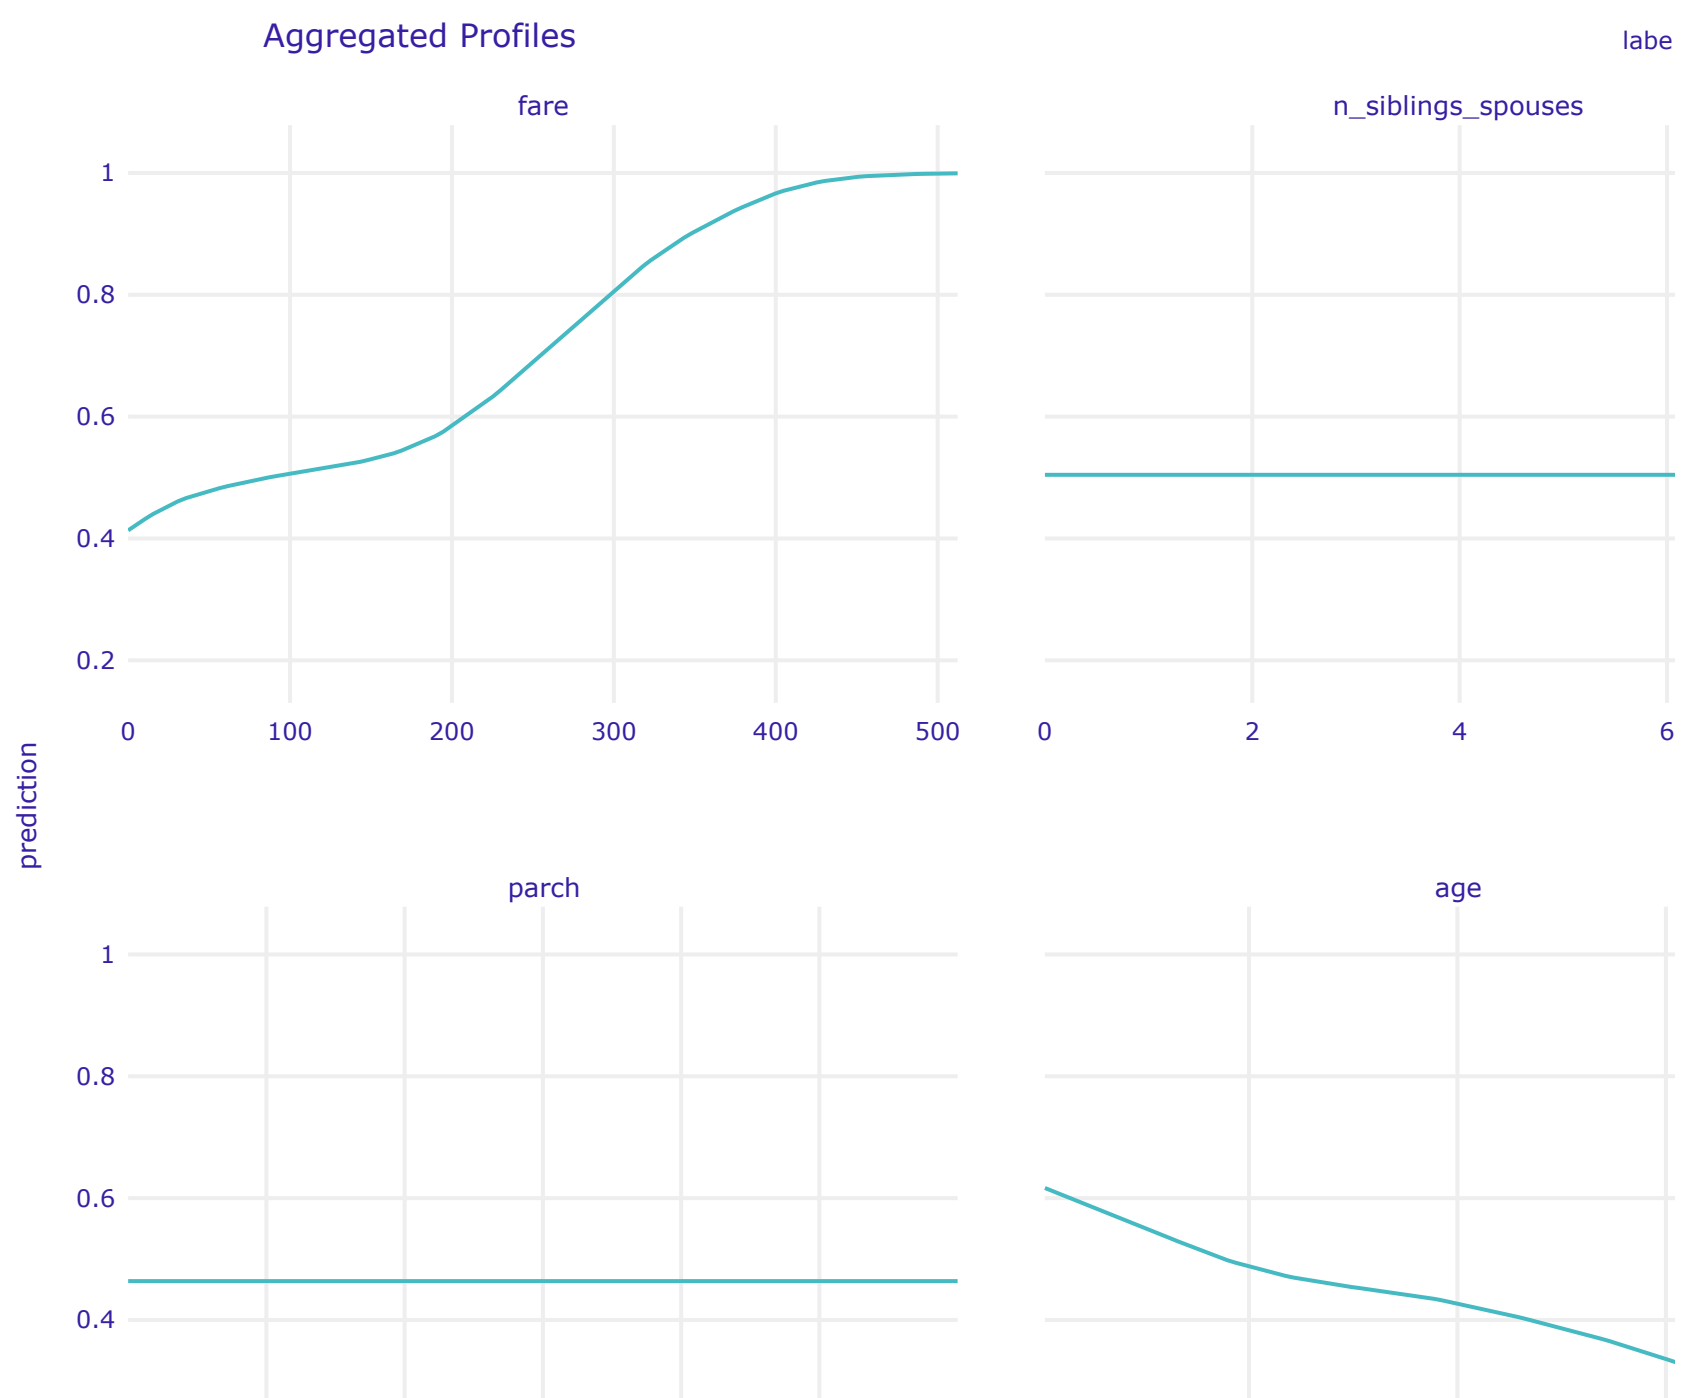
```
In [23]: explainer_keras.model_diagnostics().result
```

Out[23]:


	sex	age	n_siblings_spouses	parch	fare	class	deck	embark_town	alone	y	y_hat	residuals	abs_residuals		
0	male	22.0		1	0	7.2500	Third	unknown	Southampton	n	0	0.334225	-0.334225	0.334225	au
1	female	38.0		1	0	71.2833	First	C	Cherbourg	n	1	0.511104	0.488896	0.488896	au
2	female	26.0		0	0	7.9250	Third	unknown	Southampton	y	1	0.467287	0.532713	0.532713	au
3	female	35.0		1	0	53.1000	First	C	Southampton	n	1	0.508692	0.491308	0.491308	au
4	male	28.0		0	0	8.4583	Third	unknown	Queenstown	y	0	0.419327	-0.419327	0.419327	au
...	...	...		...	...	...	...	...	...	...	...	...	...	...	
776	female	56.0		0	1	83.1583	First	C	Cherbourg	n	1	0.498832	0.501168	0.501168	au
777	female	25.0		0	1	26.0000	Second	unknown	Southampton	n	1	0.571394	0.428606	0.428606	au
778	male	33.0		0	0	7.8958	Third	unknown	Southampton	y	0	0.369783	-0.369783	0.369783	au
779	female	39.0		0	5	29.1250	Third	unknown	Queenstown	n	0	0.501495	-0.501495	0.501495	au
780	male	26.0		0	0	30.0000	First	C	Cherbourg	y	1	0.463082	0.536918	0.536918	au

781 rows × 15 columns


```
explainer_keras.model_parts().plot()
```

[illegible][illegible]

In [26]: `explainer_keras.predict_parts(X.loc[0], type='shap').plot()`



In [27]: `explainer_keras.predict_parts(X.loc[1], type='shap').plot()`



In [28]: `X_one_hot = pd.get_dummies(X, drop_first=True)`

X\_one\_hot

Out[28]:

	age	n_siblings_spouses	parch	fare	sex_male	class_Second	class_Third	deck_B	deck_C	deck_D	deck_E	deck_F	deck_G
0	22.0	1	0	7.2500	1	0	1	0	0	0	0	0	0
1	38.0	1	0	71.2833	0	0	0	0	1	0	0	0	0
2	26.0	0	0	7.9250	0	0	1	0	0	0	0	0	0
3	35.0	1	0	53.1000	0	0	0	0	1	0	0	0	0
4	28.0	0	0	8.4583	1	0	1	0	0	0	0	0	0
...	...	...	...	...	...	...	...	...	...	...	...	...	...
776	56.0	0	1	83.1583	0	0	0	0	1	0	0	0	0
777	25.0	0	1	26.0000	0	1	0	0	0	0	0	0	0
778	33.0	0	0	7.8958	1	0	1	0	0	0	0	0	0
779	39.0	0	5	29.1250	0	0	1	0	0	0	0	0	0
780	26.0	0	0	30.0000	1	0	0	0	1	0	0	0	0

781 rows × 18 columns

In [29]: `from sklearn import tree`

```
clf = tree.DecisionTreeClassifier(max_features = 5, max_depth = 3)

X_one_hot = pd.get_dummies(X, drop_first=True)

clf = clf.fit(X_one_hot, y)
```

In [30]: `df[df.sex == 'male'].survived.value_counts()`

Out[30]:

0	382
1	106

Name: survived, dtype: int64

```
In [31]: df[df.sex == 'female'].survived.value_counts()
```

```
Out[31]: 1    217  
         0     76  
         Name: survived, dtype: int64
```

```
In [32]: df.survived.value_counts()
```

```
Out[32]: 0    458  
         1    323  
         Name: survived, dtype: int64
```

```
In [33]: #clf.classes_
```

```
In [34]: fn = list(X_one_hot.columns)

cn = ['did not survive', 'survived']

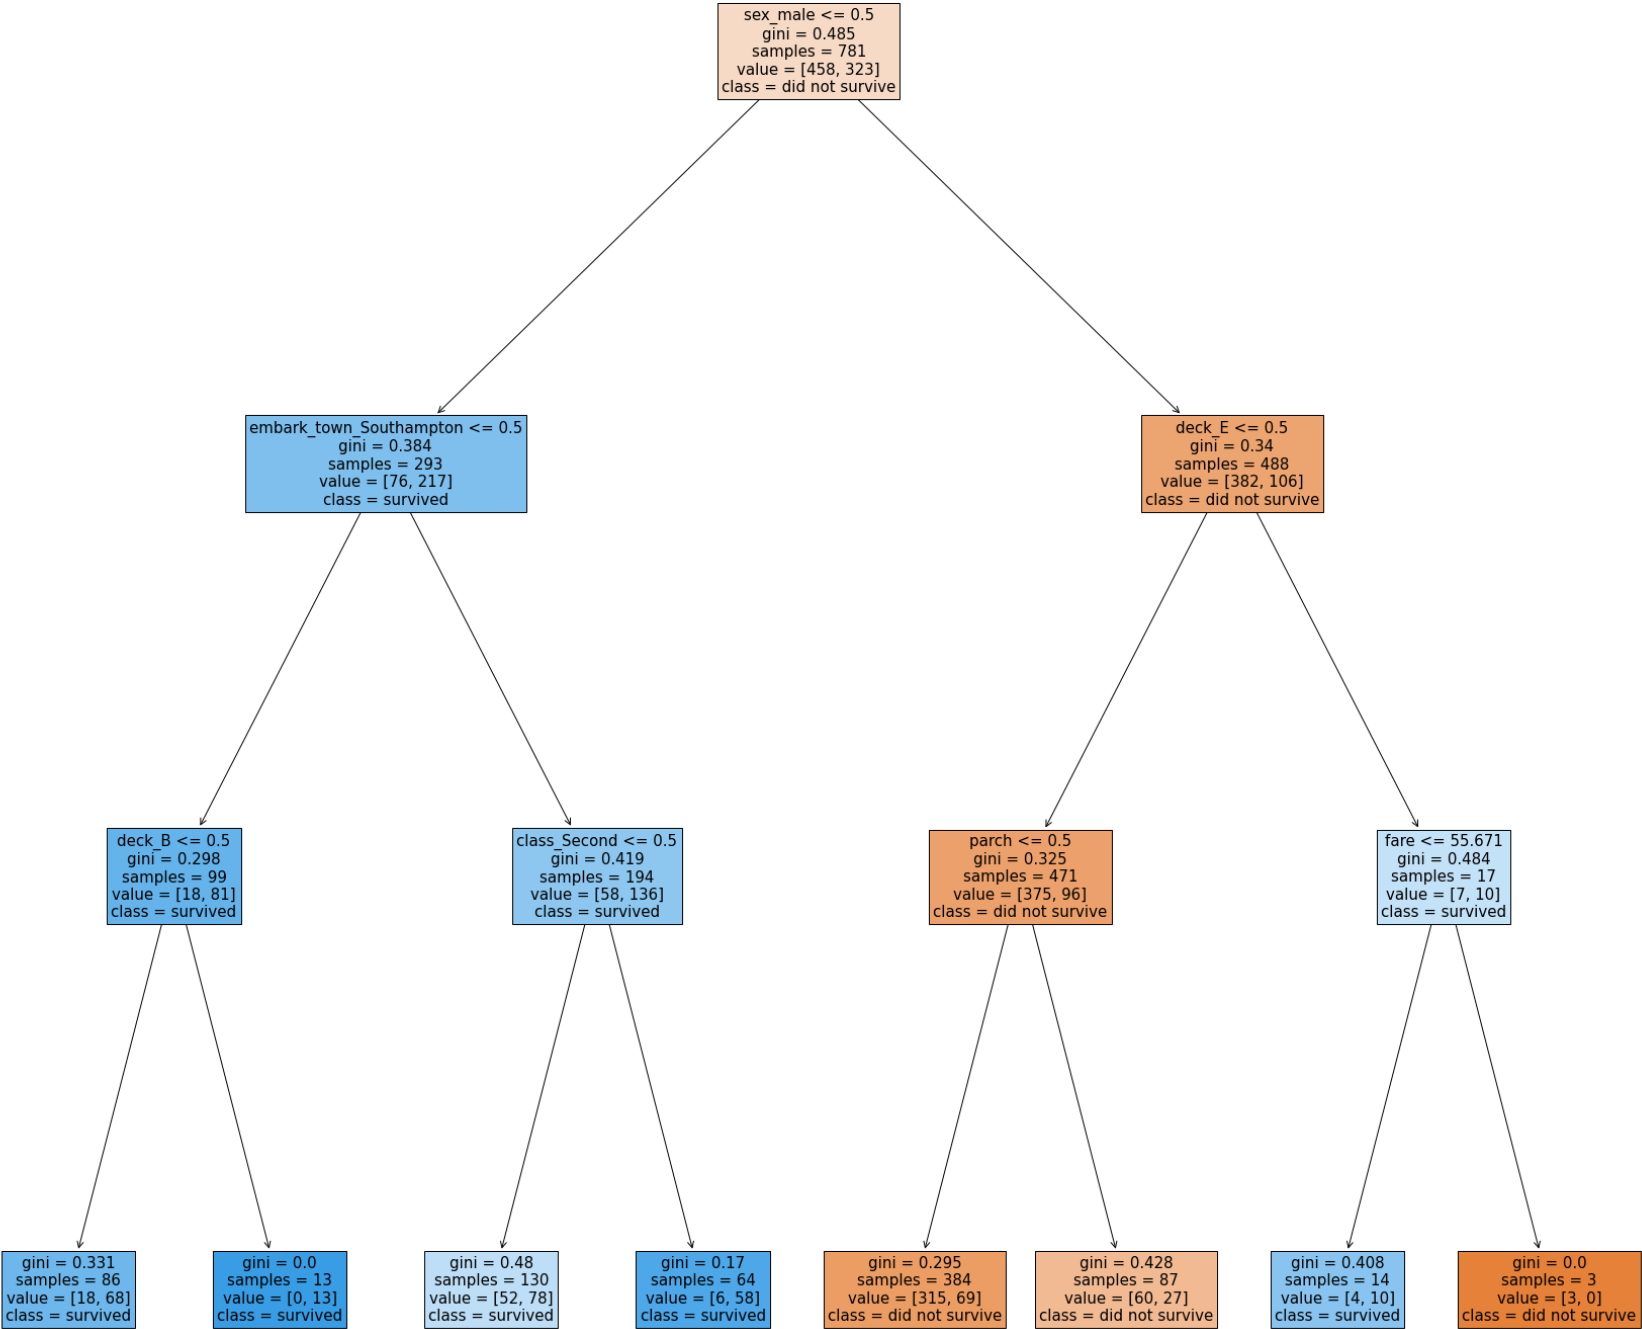
#cn = ['survived', 'did not survive']

fig, axes = plt.subplots(nrows = 1,
                          ncols = 1,
                          figsize = (30,30))
                          #dpi=500)

tree.plot_tree(clf,
               feature_names = fn,
               class_names=cn,
               filled = True, fontsize = 15)

#fig.savefig('imagename.png')
```

Out[34]: [Text(837.0, 1426.95, 'sex\_male <= 0.5\ngini = 0.485\nsamples = 781\nvalue = [458, 323]\nnclass = did not survive'),  
Text(418.5, 1019.25, 'embark\_town\_Southampton <= 0.5\ngini = 0.384\nsamples = 293\nvalue = [76, 217]\nnclass = survived'),  
Text(209.25, 611.55, 'deck\_B <= 0.5\ngini = 0.298\nsamples = 99\nvalue = [18, 81]\nnclass = survived'),  
Text(104.625, 203.84999999999999, 'gini = 0.331\nsamples = 86\nvalue = [18, 68]\nnclass = survived'),  
Text(313.875, 203.84999999999999, 'gini = 0.0\nsamples = 13\nvalue = [0, 13]\nnclass = survived'),  
Text(627.75, 611.55, 'class\_Second <= 0.5\ngini = 0.419\nsamples = 194\nvalue = [58, 136]\nnclass = survived'),  
Text(523.125, 203.84999999999999, 'gini = 0.48\nsamples = 130\nvalue = [52, 78]\nnclass = survived'),  
Text(732.375, 203.84999999999999, 'gini = 0.17\nsamples = 64\nvalue = [6, 58]\nnclass = survived'),  
Text(1255.5, 1019.25, 'deck\_E <= 0.5\ngini = 0.34\nsamples = 488\nvalue = [382, 106]\nnclass = did not survive'),  
Text(1046.25, 611.55, 'parch <= 0.5\ngini = 0.325\nsamples = 471\nvalue = [375, 96]\nnclass = did not survive'),  
Text(941.625, 203.84999999999999, 'gini = 0.295\nsamples = 384\nvalue = [315, 69]\nnclass = did not survive'),  
Text(1150.875, 203.84999999999999, 'gini = 0.428\nsamples = 87\nvalue = [60, 27]\nnclass = did not survive'),  
Text(1464.75, 611.55, 'fare <= 55.671\ngini = 0.484\nsamples = 17\nvalue = [7, 10]\nnclass = survived'),  
Text(1360.125, 203.84999999999999, 'gini = 0.408\nsamples = 14\nvalue = [4, 10]\nnclass = survived'),  
Text(1569.375, 203.84999999999999, 'gini = 0.0\nsamples = 3\nvalue = [3, 0]\nnclass = did not survive')]



working version

```
In [35]: import dalex as dx

data = pd.read_csv("https://raw.githubusercontent.com/pbiecek/xai-happiness/main/happiness.csv", index_col=0)
data.head()
```

Out[35]:

	score	gdp_per_capita	social_support	healthy_life_expectancy	freedom_to_make_life_choices	generosity	perceptions_of_cor
Afghanistan	3.203	0.350	0.517	0.361	0.000	0.158	
Albania	4.719	0.947	0.848	0.874	0.383	0.178	
Algeria	5.211	1.002	1.160	0.785	0.086	0.073	
Argentina	6.086	1.092	1.432	0.881	0.471	0.066	
Armenia	4.559	0.850	1.055	0.815	0.283	0.095	

```
In [36]: X, y = data.drop('score', axis=1), data.score
n, p = X.shape

X
```

Out[36]:

	gdp_per_capita	social_support	healthy_life_expectancy	freedom_to_make_life_choices	generosity	perceptions_of_corruption
Afghanistan	0.350	0.517	0.361	0.000	0.158	0.025
Albania	0.947	0.848	0.874	0.383	0.178	0.027
Algeria	1.002	1.160	0.785	0.086	0.073	0.114
Argentina	1.092	1.432	0.881	0.471	0.066	0.050
Armenia	0.850	1.055	0.815	0.283	0.095	0.064
...	...	...	...	...	...	...
Venezuela	0.960	1.427	0.805	0.154	0.064	0.047
Vietnam	0.741	1.346	0.851	0.543	0.147	0.073
Yemen	0.287	1.163	0.463	0.143	0.108	0.077
Zambia	0.578	1.058	0.426	0.431	0.247	0.087
Zimbabwe	0.366	1.114	0.433	0.361	0.151	0.089

156 rows × 6 columns

```
In [37]: y
```

Out[37]: Afghanistan 3.203  
Albania 4.719  
Algeria 5.211  
Argentina 6.086  
Armenia 4.559  
...  
Venezuela 4.707  
Vietnam 5.175  
Yemen 3.380  
Zambia 4.107  
Zimbabwe 3.663  
Name: score, Length: 156, dtype: float64

```
In [38]: #tf.random.set_seed(11)

normalizer = tf.keras.layers.experimental.preprocessing.Normalization(input_shape=[p,])
normalizer.adapt(X.to_numpy())

model = tf.keras.Sequential([
    normalizer,
    tf.keras.Input(shape=(p,)),
    tf.keras.layers.Dense(p*2, activation='relu'),
    tf.keras.layers.Dense(p*3, activation='relu'),
    tf.keras.layers.Dense(p*2, activation='relu'),
    tf.keras.layers.Dense(p, activation='relu'),
    tf.keras.layers.Dense(1, activation='linear')
])

model.compile(
    optimizer=tf.keras.optimizers.Adam(0.001),
    loss=tf.keras.losses.mae
)

model.fit(X, y, batch_size=int(n/10), epochs=2000, verbose=False)
```

WARNING:tensorflow:Please add `keras.layers.InputLayer` instead of `keras.Input` to Sequential model. `keras.Input` is intended to be used by Functional model.

Out[38]: <tensorflow.python.keras.callbacks.History at 0x1d3a2287340>

```
In [39]: #type(model)
```

```
In [40]: #model.output_shape
```

```
In [41]: explainer = dx.Explainer(model, X, y, label='happiness')
```

Preparation of a new explainer is initiated

```
-> data          : 156 rows 6 cols
-> target variable : Parameter 'y' was a pandas.Series. Converted to a numpy.ndarray.
-> target variable : 156 values
-> model_class    : tensorflow.python.keras.engine.sequential.Sequential (default)
-> label          : happiness
-> predict function : <function yhat_tf_regression at 0x000001D1C4457790> will be used (default)
-> predict function : Accepts pandas.DataFrame and numpy.ndarray.
-> predicted values : min = 2.88, mean = 5.4, max = 7.56
-> model type      : regression will be used (default)
-> residual function : difference between y and yhat (default)
-> residuals       : min = -0.63, mean = 0.0104, max = 0.697
-> model_info      : package tensorflow
```

A new explainer has been created!

```
In [42]: #explainer_new = dx.Explainer(model, X, y, label='happiness',
#                                     predict_function = dx._explainer.yhat.yhat_tf_regression)
```

```
In [43]: #explainer.predict_function
```

```
In [44]: #dx._explainer.yhat.yhat_tf_regression(model, X)
```

```
In [45]: #explainer.residual_function
```

```
In [46]: #dx._explainer.checks.check_residual_function.residual_function

#def rf(_model, _data, _y):
#    return _y - dx._explainer.yhat.yhat_tf_regression(model, X)

#rf(model, X, y)
```


```
In [47]: explainer.model_performance()
```

Out[47]:

	mse	rmse	r2	mae	mad
happiness	0.025321	0.159127	0.979432	0.100341	0.05421

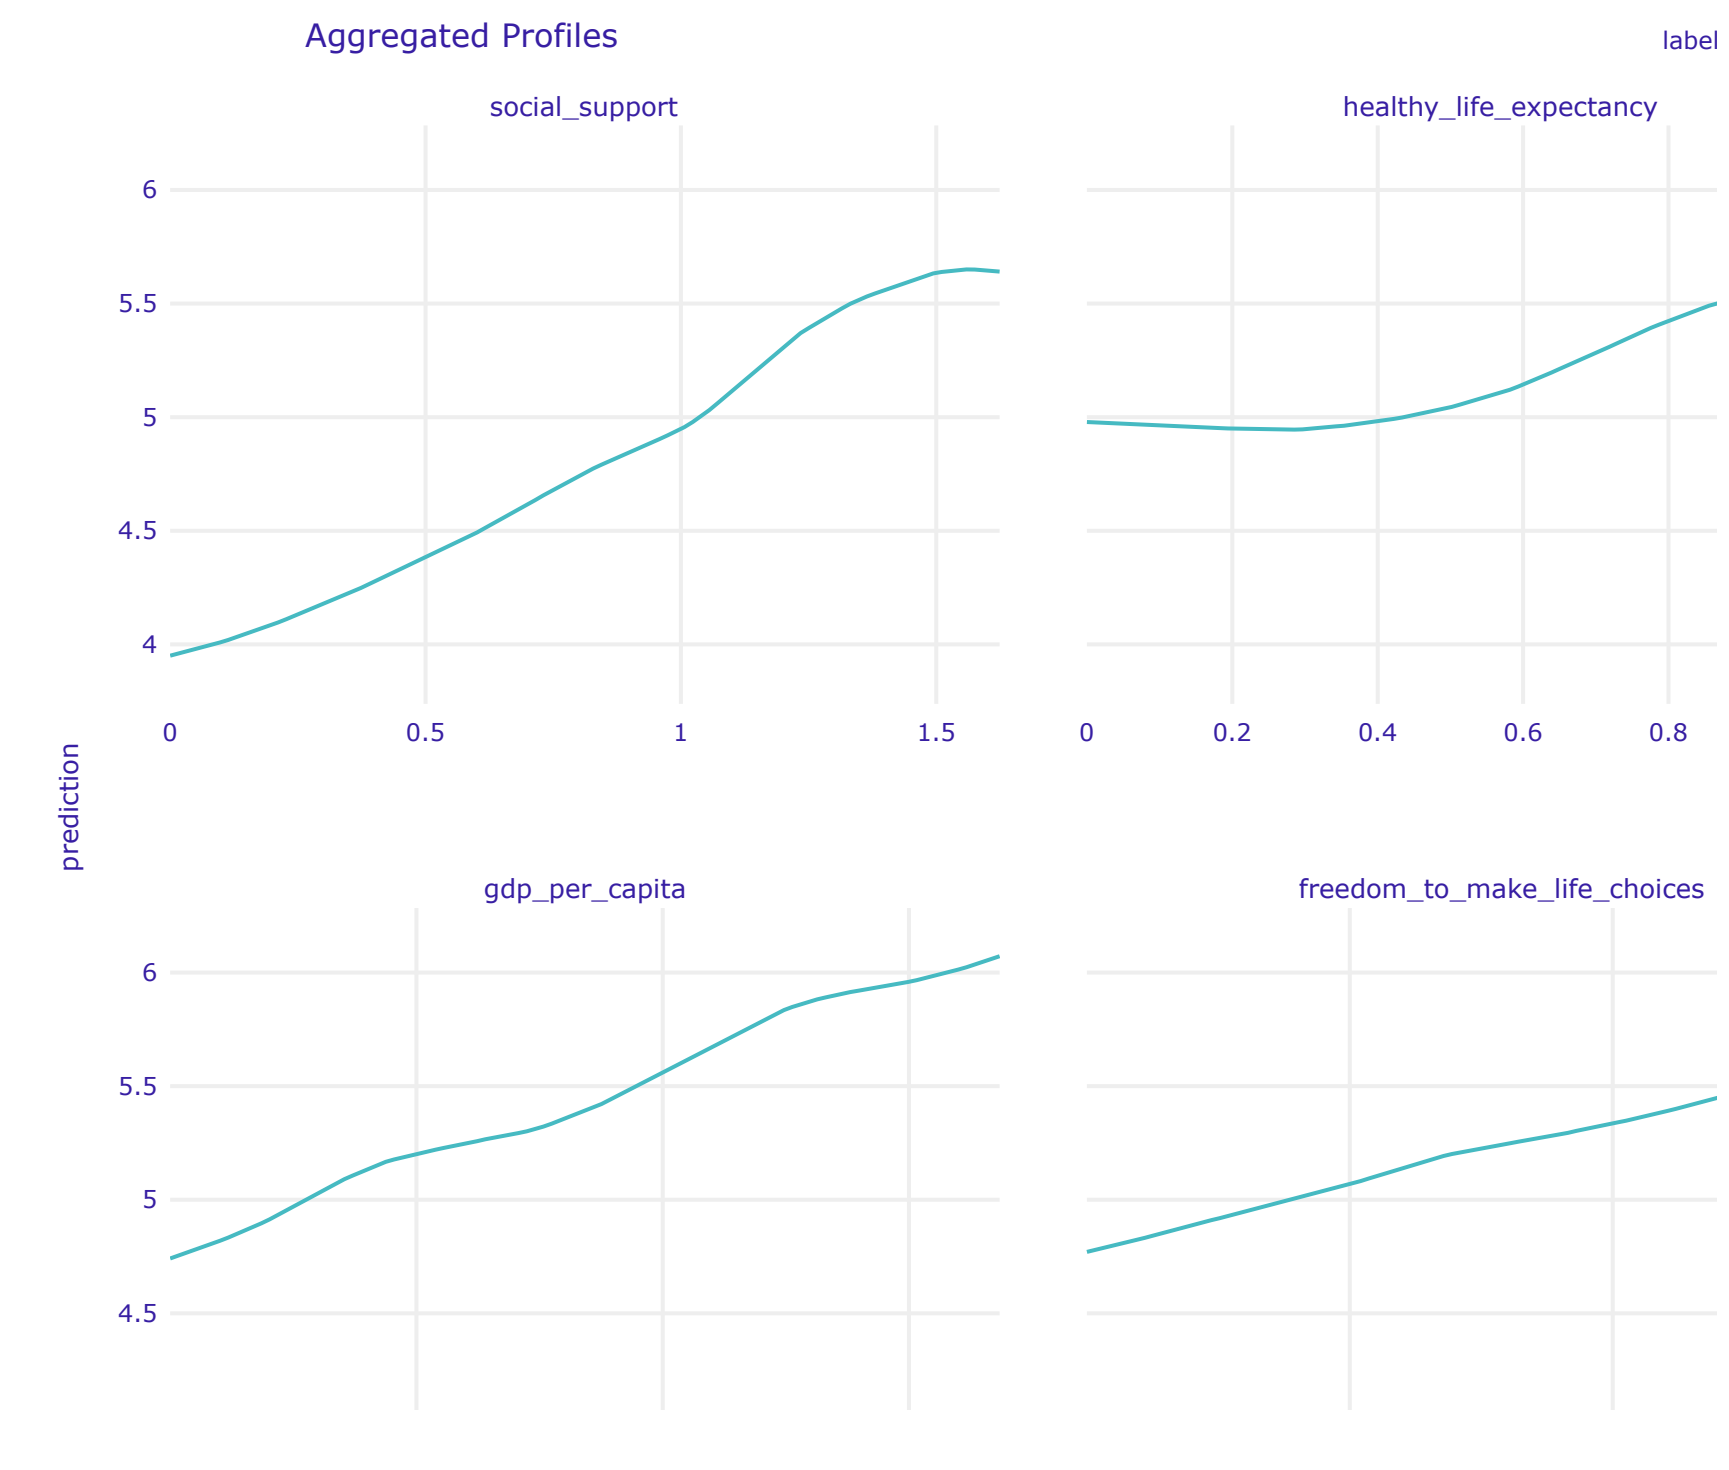


```
explainer.model_parts().plot()
```




```
explainer.model_profile().plot(variables=['social_support', 'healthy_life_expectancy',  
                                         'gdp_per_capita', 'freedom_to_make_life_choices'])
```

```
Calculating ceteris paribus: 100%|███████████| 6/6 [00:03<00:00,  
1.99it/s]
```



```
In [50]: explainer.model_diagnostics().plot(variable='social_support', yvariable="abs_residuals", marker_size=5, line_wid
```



```
In [51]: explainer.model_diagnostics().result
```

Out[51]:

	gdp_per_capita	social_support	healthy_life_expectancy	freedom_to_make_life_choices	generosity	perceptions_of_corruption
Afghanistan	0.350	0.517	0.361	0.000	0.158	0.025
Albania	0.947	0.848	0.874	0.383	0.178	0.027
Algeria	1.002	1.160	0.785	0.086	0.073	0.114
Argentina	1.092	1.432	0.881	0.471	0.066	0.050
Armenia	0.850	1.055	0.815	0.283	0.095	0.064
...	...	...	...	...	...	...
Venezuela	0.960	1.427	0.805	0.154	0.064	0.047
Vietnam	0.741	1.346	0.851	0.543	0.147	0.073
Yemen	0.287	1.163	0.463	0.143	0.108	0.077
Zambia	0.578	1.058	0.426	0.431	0.247	0.087
Zimbabwe	0.366	1.114	0.433	0.361	0.151	0.089


156 rows × 12 columns

```
In [52]: explainer.predict_parts(X.loc['Poland'], type='shap').plot()
```

```
Exception ignored in: <function CapturableResource.__del__ at 0x000001D1E4C61EE0>
Traceback (most recent call last):
  File "C:\Users\deepl\anaconda3\lib\site-packages\tensorflow\python\training\tracking\tracking.py", line 27
7, in __del__
    self._destroy_resource()
  File "C:\Users\deepl\anaconda3\lib\site-packages\tensorflow\python\eager\def_function.py", line 889, in __
call__
    result = self._call(*args, **kwargs)
  File "C:\Users\deepl\anaconda3\lib\site-packages\tensorflow\python\eager\def_function.py", line 924, in _c
all
    results = self._stateful_fn(*args, **kwargs)
  File "C:\Users\deepl\anaconda3\lib\site-packages\tensorflow\python\eager\function.py", line 3022, in __cal
l__
    filtered_flat_args) = self._maybe_define_function(args, kwargs)
  File "C:\Users\deepl\anaconda3\lib\site-packages\tensorflow\python\eager\function.py", line 3444, in _mayb
e_define_function
    graph_function = self._create_graph_function(args, kwargs)
  File "C:\Users\deepl\anaconda3\lib\site-packages\tensorflow\python\eager\function.py", line 3279, in _crea
te_graph_function
    func_graph_module.func_graph_from_py_func(
  File "C:\Users\deepl\anaconda3\lib\site-packages\tensorflow\python\framework\func_graph.py", line 999, in
func_graph_from_py_func
    func_outputs = python_func(*func_args, **func_kwargs)
  File "C:\Users\deepl\anaconda3\lib\site-packages\tensorflow\python\eager\def_function.py", line 672, in wr
apped_fn
    out = weak_wrapped_fn().__wrapped__(*args, **kwargs)
AttributeError: 'NoneType' object has no attribute '__wrapped__'
Exception ignored in: <function CapturableResource.__del__ at 0x000001D1E4C61EE0>
Traceback (most recent call last):
  File "C:\Users\deepl\anaconda3\lib\site-packages\tensorflow\python\training\tracking\tracking.py", line 27
7, in __del__
    self._destroy_resource()
  File "C:\Users\deepl\anaconda3\lib\site-packages\tensorflow\python\eager\def_function.py", line 889, in __
call__
    result = self._call(*args, **kwargs)
  File "C:\Users\deepl\anaconda3\lib\site-packages\tensorflow\python\eager\def_function.py", line 924, in _c
all
    results = self._stateful_fn(*args, **kwargs)
  File "C:\Users\deepl\anaconda3\lib\site-packages\tensorflow\python\eager\function.py", line 3022, in __cal
l__
    filtered_flat_args) = self._maybe_define_function(args, kwargs)
  File "C:\Users\deepl\anaconda3\lib\site-packages\tensorflow\python\eager\function.py", line 3444, in _mayb
e_define_function
    graph_function = self._create_graph_function(args, kwargs)
  File "C:\Users\deepl\anaconda3\lib\site-packages\tensorflow\python\eager\function.py", line 3279, in _crea
te_graph_function
    func_graph_module.func_graph_from_py_func(
  File "C:\Users\deepl\anaconda3\lib\site-packages\tensorflow\python\framework\func_graph.py", line 999, in
func_graph_from_py_func
    func_outputs = python_func(*func_args, **func_kwargs)
  File "C:\Users\deepl\anaconda3\lib\site-packages\tensorflow\python\eager\def_function.py", line 672, in wr
apped_fn
    out = weak_wrapped_fn().__wrapped__(*args, **kwargs)
AttributeError: 'NoneType' object has no attribute '__wrapped__'
Exception ignored in: <function CapturableResource.__del__ at 0x000001D1E4C61EE0>
Traceback (most recent call last):
  File "C:\Users\deepl\anaconda3\lib\site-packages\tensorflow\python\training\tracking\tracking.py", line 27
7, in __del__
    self._destroy_resource()
  File "C:\Users\deepl\anaconda3\lib\site-packages\tensorflow\python\eager\def_function.py", line 889, in __
call__
    result = self._call(*args, **kwargs)
  File "C:\Users\deepl\anaconda3\lib\site-packages\tensorflow\python\eager\def_function.py", line 924, in _c
all
    results = self._stateful_fn(*args, **kwargs)
  File "C:\Users\deepl\anaconda3\lib\site-packages\tensorflow\python\eager\function.py", line 3022, in __cal
l__
    filtered_flat_args) = self._maybe_define_function(args, kwargs)
  File "C:\Users\deepl\anaconda3\lib\site-packages\tensorflow\python\eager\function.py", line 3444, in _mayb
e_define_function
    graph_function = self._create_graph_function(args, kwargs)
  File "C:\Users\deepl\anaconda3\lib\site-packages\tensorflow\python\eager\function.py", line 3279, in _crea
te_graph_function
    func_graph_module.func_graph_from_py_func(
  File "C:\Users\deepl\anaconda3\lib\site-packages\tensorflow\python\framework\func_graph.py", line 999, in
func_graph_from_py_func
    func_outputs = python_func(*func_args, **func_kwargs)
  File "C:\Users\deepl\anaconda3\lib\site-packages\tensorflow\python\eager\def_function.py", line 672, in wr
apped_fn
    out = weak_wrapped_fn().__wrapped__(*args, **kwargs)
AttributeError: 'NoneType' object has no attribute '__wrapped__'
Exception ignored in: <function CapturableResource.__del__ at 0x000001D1E4C61EE0>
Traceback (most recent call last):
  File "C:\Users\deepl\anaconda3\lib\site-packages\tensorflow\python\training\tracking\tracking.py", line 27
7, in __del__
    self._destroy_resource()
  File "C:\Users\deepl\anaconda3\lib\site-packages\tensorflow\python\eager\def_function.py", line 889, in __
```

```
call__
    result = self._call(*args, **kwargs)
    File "C:\Users\deepl\anaconda3\lib\site-packages\tensorflow\python\eager\def_function.py", line 924, in _c
all
    results = self._stateful_fn(*args, **kwargs)
    File "C:\Users\deepl\anaconda3\lib\site-packages\tensorflow\python\eager\function.py", line 3022, in __cal
l__
    filtered_flat_args) = self._maybe_define_function(args, kwargs)
    File "C:\Users\deepl\anaconda3\lib\site-packages\tensorflow\python\eager\function.py", line 3444, in _mayb
e_define_function
    graph_function = self._create_graph_function(args, kwargs)
    File "C:\Users\deepl\anaconda3\lib\site-packages\tensorflow\python\eager\function.py", line 3279, in _crea
te_graph_function
    func_graph_module.func_graph_from_py_func(
    File "C:\Users\deepl\anaconda3\lib\site-packages\tensorflow\python\framework\func_graph.py", line 999, in
func_graph_from_py_func
    func_outputs = python_func(*func_args, **func_kwargs)
    File "C:\Users\deepl\anaconda3\lib\site-packages\tensorflow\python\eager\def_function.py", line 672, in wr
apped_fn
    out = weak_wrapped_fn().__wrapped__(*args, **kwargs)
AttributeError: 'NoneType' object has no attribute '__wrapped__'
Exception ignored in: <function CapturableResource.__del__ at 0x000001D1E4C61EE0>
Traceback (most recent call last):
  File "C:\Users\deepl\anaconda3\lib\site-packages\tensorflow\python\training\tracking\tracking.py", line 27
7, in __del__
    self._destroy_resource()
  File "C:\Users\deepl\anaconda3\lib\site-packages\tensorflow\python\eager\def_function.py", line 889, in __
call__
    result = self._call(*args, **kwargs)
    File "C:\Users\deepl\anaconda3\lib\site-packages\tensorflow\python\eager\def_function.py", line 924, in _c
all
    results = self._stateful_fn(*args, **kwargs)
    File "C:\Users\deepl\anaconda3\lib\site-packages\tensorflow\python\eager\function.py", line 3022, in __cal
l__
    filtered_flat_args) = self._maybe_define_function(args, kwargs)
    File "C:\Users\deepl\anaconda3\lib\site-packages\tensorflow\python\eager\function.py", line 3444, in _mayb
e_define_function
    graph_function = self._create_graph_function(args, kwargs)
    File "C:\Users\deepl\anaconda3\lib\site-packages\tensorflow\python\eager\function.py", line 3279, in _crea
te_graph_function
    func_graph_module.func_graph_from_py_func(
    File "C:\Users\deepl\anaconda3\lib\site-packages\tensorflow\python\framework\func_graph.py", line 999, in
func_graph_from_py_func
    func_outputs = python_func(*func_args, **func_kwargs)
    File "C:\Users\deepl\anaconda3\lib\site-packages\tensorflow\python\eager\def_function.py", line 672, in wr
apped_fn
    out = weak_wrapped_fn().__wrapped__(*args, **kwargs)
AttributeError: 'NoneType' object has no attribute '__wrapped__'
Exception ignored in: <function CapturableResource.__del__ at 0x000001D1E4C61EE0>
Traceback (most recent call last):
  File "C:\Users\deepl\anaconda3\lib\site-packages\tensorflow\python\training\tracking\tracking.py", line 27
7, in __del__
    self._destroy_resource()
  File "C:\Users\deepl\anaconda3\lib\site-packages\tensorflow\python\eager\def_function.py", line 889, in __
call__
    result = self._call(*args, **kwargs)
    File "C:\Users\deepl\anaconda3\lib\site-packages\tensorflow\python\eager\def_function.py", line 924, in _c
all
    results = self._stateful_fn(*args, **kwargs)
    File "C:\Users\deepl\anaconda3\lib\site-packages\tensorflow\python\eager\function.py", line 3022, in __cal
l__
    filtered_flat_args) = self._maybe_define_function(args, kwargs)
    File "C:\Users\deepl\anaconda3\lib\site-packages\tensorflow\python\eager\function.py", line 3444, in _mayb
```

```
e_define_function
    graph_function = self._create_graph_function(args, kwargs)
    File "C:\Users\deep1\anaconda3\lib\site-packages\tensorflow\python\eager\function.py", line 3279, in _crea
te_graph_function
    func_graph_module.func_graph_from_py_func(
    File "C:\Users\deep1\anaconda3\lib\site-packages\tensorflow\python\framework\func_graph.py", line 999, in
func_graph_from_py_func
    func_outputs = python_func(*func_args, **func_kwargs)
    File "C:\Users\deep1\anaconda3\lib\site-packages\tensorflow\python\eager\def_function.py", line 672, in wr
apped_fn
    out = weak_wrapped_fn().__wrapped__(*args, **kws)
AttributeError: 'NoneType' object has no attribute '__wrapped__'
```




```
In [53]: pp_list = []

for country in ['Afghanistan', 'Belgium', 'China', 'Denmark', 'Ethiopia']:
    pp = explainer.predict_parts(X.loc[country], type='break_down')
    pp.result.label = country
    pp_list += [pp]


pp_list[0].plot(pp_list[1::], min_max=[2.5, 8.5])
```

Break Down



```
In [54]: lime_explanation = explainer.predict_surrogate(X.loc['United States'], mode='regression')

lime_explanation.plot()
```



```
In [55]: lime_explanation.result
```

Out[55]:

	variable	effect
0	gdp_per_capita > 1.23	0.553652
1	social_support > 1.45	0.467272
2	0.09 < perceptions_of_corruption <= 0.14	-0.264506
3	generosity > 0.25	-0.087135
4	0.79 < healthy_life_expectancy <= 0.88	0.065358
5	0.42 < freedom_to_make_life_choices <= 0.51	0.045682

```
In [56]: surrogate_model = explainer.model_surrogate(max_vars=4, max_depth=3)

surrogate_model.performance
```

Out[56]:

	mse	rmse	r2	mae	mad
DecisionTreeRegressor	0.195621	0.442291	0.820771	0.353305	0.286188

```
In [57]: surrogate_model.plot()
```

