

Real-World Deployment of ML in Production Systems

Author: Hussain Abbas, MSc

© 2021 Stats AI LLC

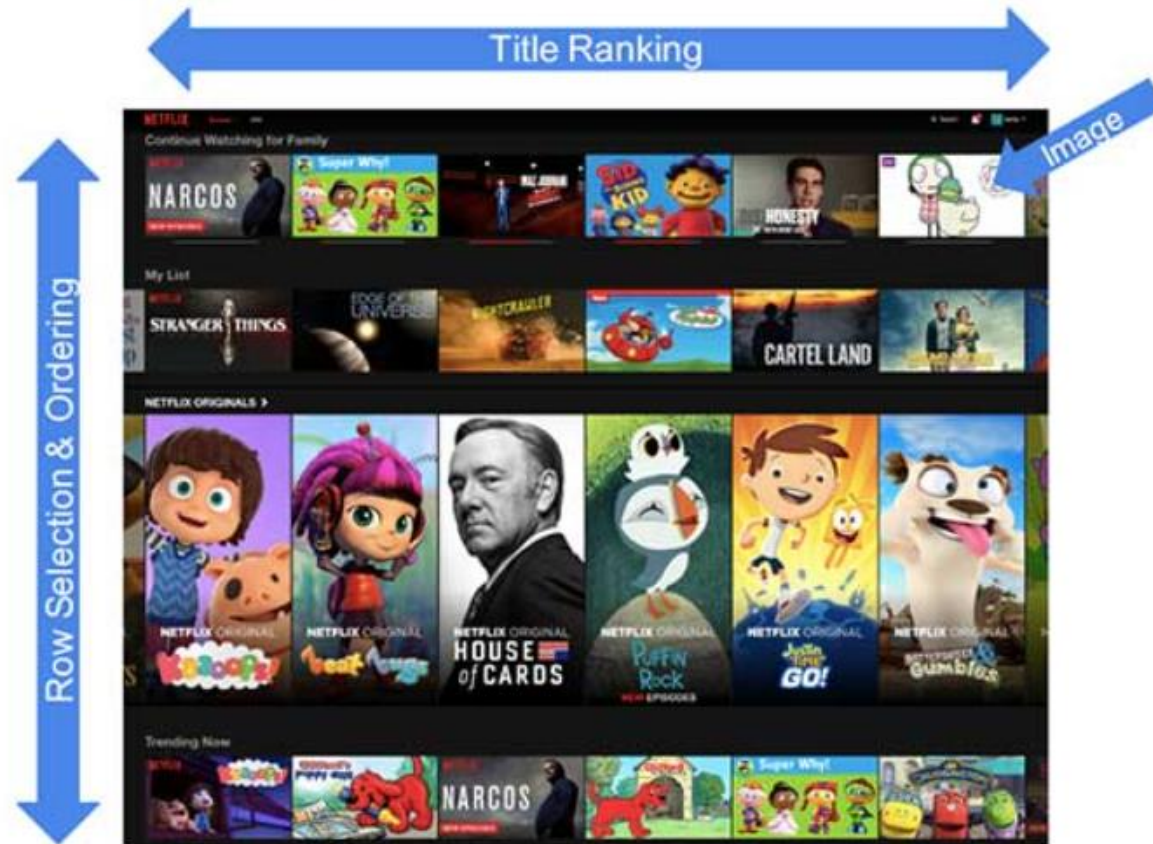
All Rights Reserved

Why Deploy ML Models?

- ML model Deployment enables firms to create new services!
- The interaction between the user and the ML model is the service
 - Recommender Systems: **Netflix Recommendation Algorithm**
 - Financial Services: Algorithmic Portfolio Selection
 - Rideshare: Uber arrival time prediction
- Deployed ML models enable firms to:
 - Acquire new customers
 - Retain existing customers
 - Grow revenues
 - Cut costs
 - Gain market share
 - Gain competitive advantage
 - Stay relevant

Recommender Systems

Everything is a Recommendation



NETFLIX

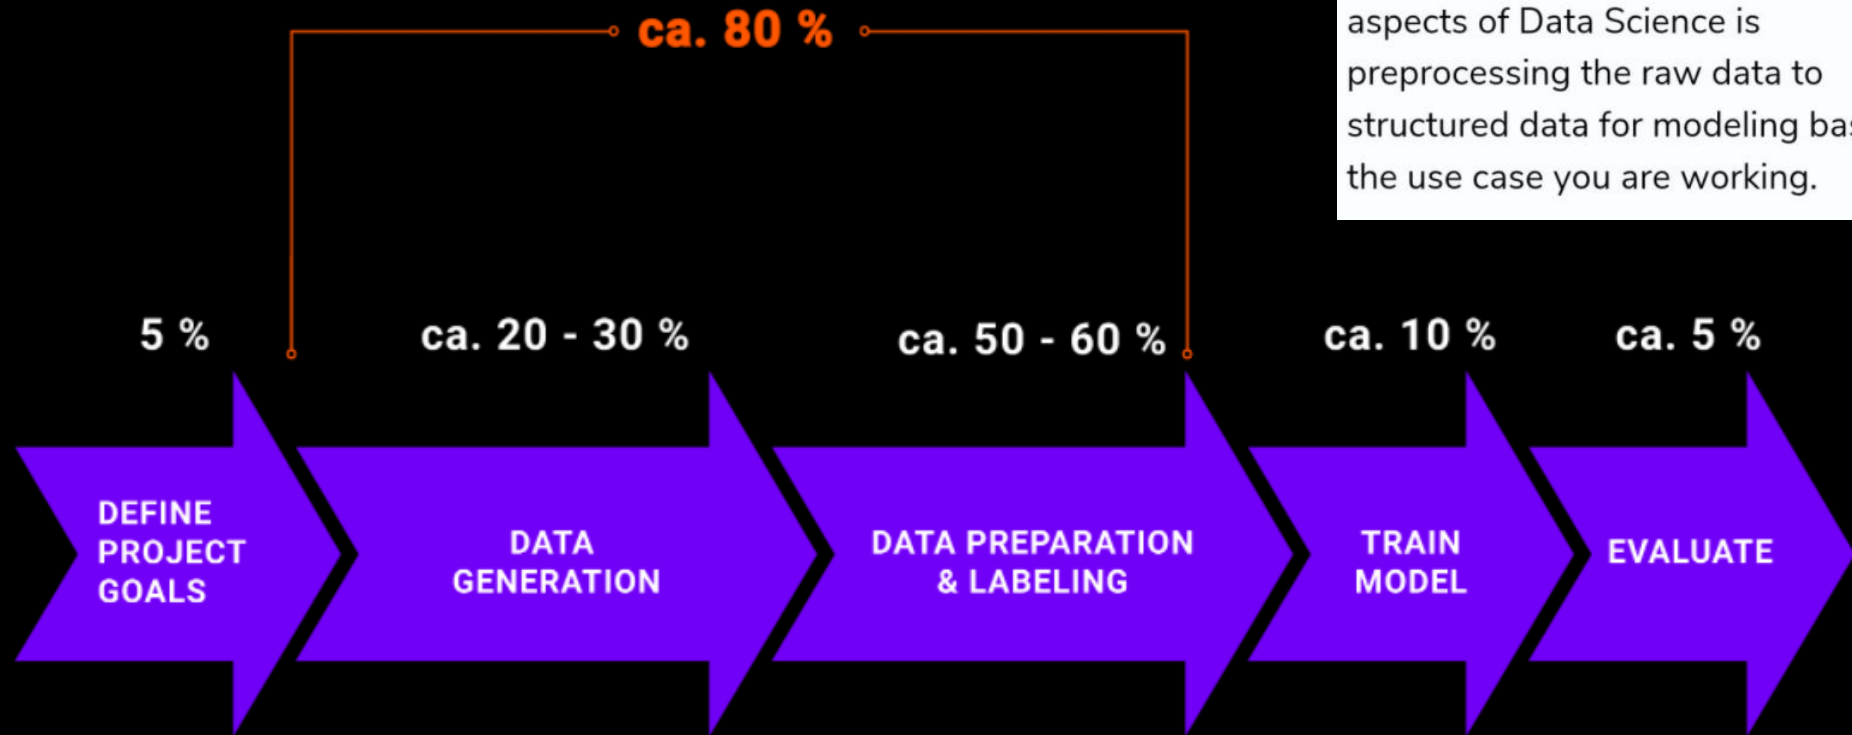
Recommendations are driven by machine learning algorithms

Over 80% of what members watch comes from our recommendations

Stats AI

TYPICAL MACHINE LEARNING PROJECT

Phases and estimated time in percent



DATA GENERATION

- Acquire data (search, make or buy)
- Data generation
- Data Augmentation

DATA PREPARATION

- Store / load data
- Organize data
- Correct, normalize ..
- Label & annotate data

TRAIN & EVALUATE MODEL

- Choose model
- Train model
- Evaluate model
- Deploy model





You are here!

Model Deployment

- Once an ML model is developed, it needs to be deployed as a REST API so users can interact with it
- How a ML REST API works under the hood:
 - The trained ML model is pickled
 - JSON payload delivers feature data via a POST request
 - Various transformations are applied to the JSON payload before it enters the model
 - The model takes in the input and an output prediction is generated
 - The API returns the prediction from the model in JSON
 - The JSON output is presented to the user either directly or indirectly
 - Direct: Streamlit, Gradio, GUI, etc.
 - Indirect: API output is an input into an Enterprise application
- The two types of deployment
 - Batch
 - Take in a set of inputs at a point in time, and generate predictions
 - Real-time
 - Generate predictions on the fly as requested

Model Deployment

- Where does the JSON payload that triggers the POST Request come from?
 - A database
 - User interaction
 - Manual user entry
 - User Location
 - Streaming data
 - All the above combined!

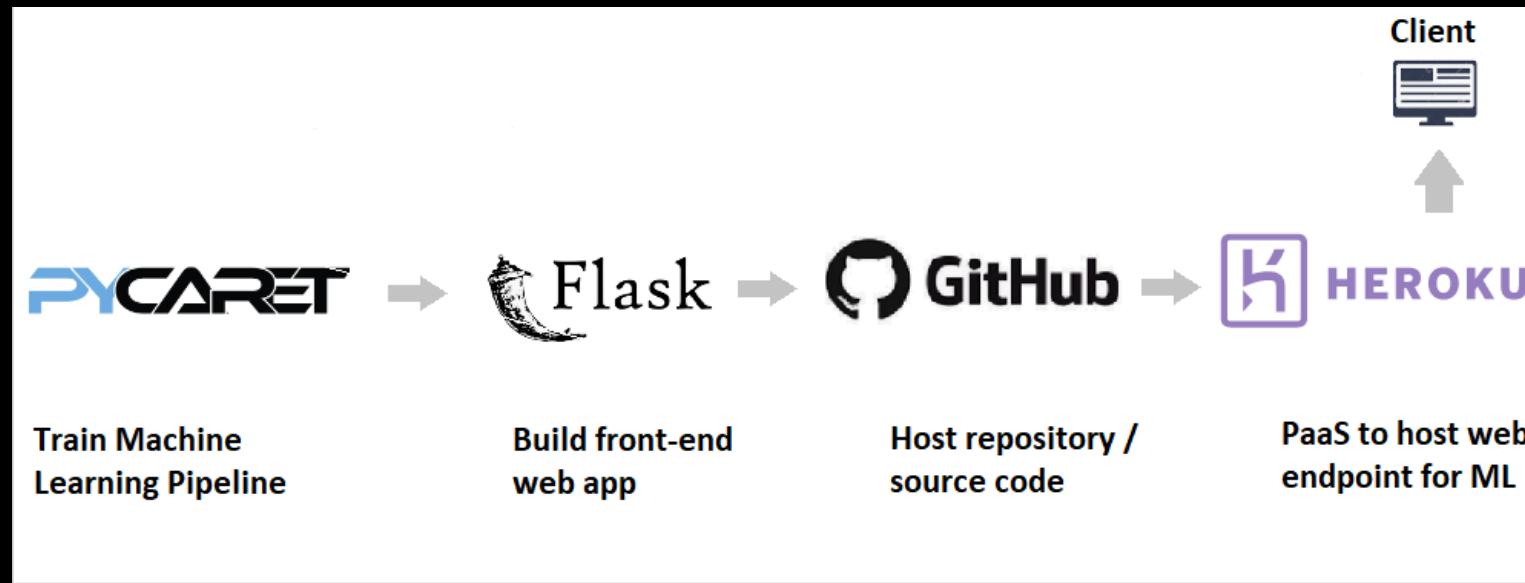


Model Deployment

- Generally, external customer facing services require a level of engineering effort magnitudes larger than internal services
- As internal services become more complex, they start to approximate external services
- Thus, what is needed are streamlined deployment systems that enable rapid iteration that can scale with complexity and demand ... hence the Cloud

Deployment in the Cloud

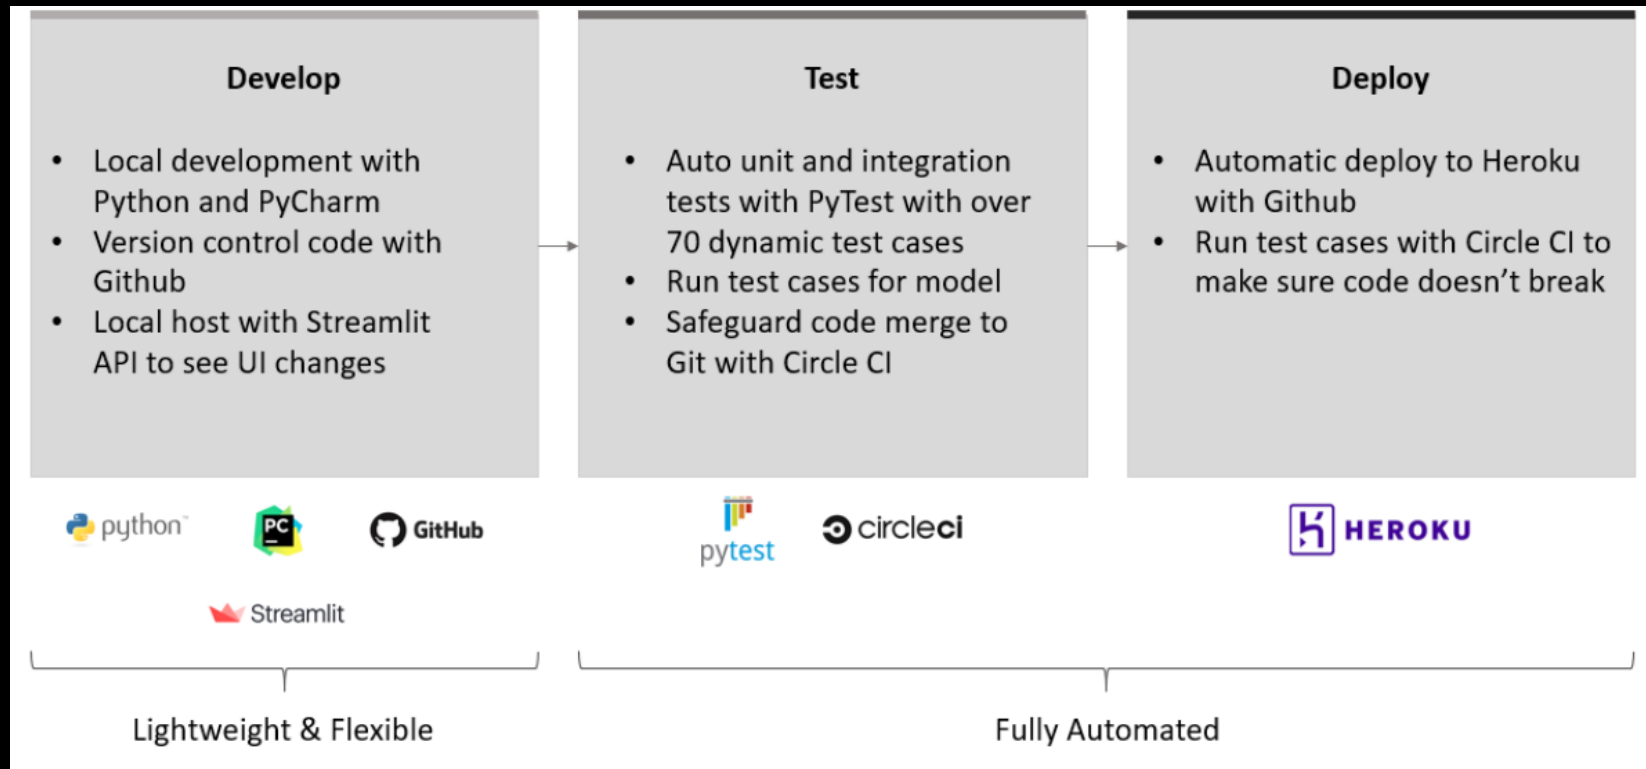
Example Workflow: Level I



<https://towardsdatascience.com/build-and-deploy-your-first-machine-learning-web-app-e020db344a99>

- Pycaret: Trains ML model and pickles saved model
- HTML: Website the user interacts with (in this case a simple web form)
- Flask: Backend web framework which powers the REST API and HTML front-end
- GitHub: Where all the code necessary to create and deploy the model lives
- Heroku: PAAS (Platform as a Service) used to quickly develop and host apps

Example Workflow: Level II



<https://www.kdnuggets.com/2020/02/machine-learning-challenge-build-deploy-app-streamlit-devops.html>


- Unit testing: ensures that an individual module behaves as expected
- Integration testing: ensure that a collection of modules interoperate as expected
- Circle CI: gate check code merge on Github as well as deployment of ML models to production

Docker & Kubernetes

- Container-based applications can be moved easily from on-prem systems to cloud environments or from developers' laptops to servers
- Dockerfile
 - a text file with the instructions to build a Docker image
 - specifies the OS, the programming languages, environmental variables, file locations, network ports, and other components the application needs
- Docker Image:
 - is a portable file containing the actual files. Thus, it is usually a large file
 - can be stacked. i.e., you can take an image someone else created and build upon it
- Docker Container: an instance of the docker image, i.e., the program that you want to run
- Container Registry: DockerHub, Azure Container Registry Amazon ECR, Google Container Registry
 - Where we upload Docker images we create and download docker images created by others
- Kubernetes: Enables container orchestration, i.e, the ability for containers to talk to each other
- We can create multiple containers from a single image on a single host OS
- Kubernetes is how we enable those containers to talk to each other across multiple host OS
- Thus, scaling up an application simply becomes spinning up more machines in a Kubernetes cluster

Example Workflow: Level III

Containerize and Deploy Machine Learning Pipeline as an Azure Web App




<https://towardsdatascience.com/deploy-machine-learning-pipeline-on-cloud-using-docker-container-bec64458dc01>

Stats AI

- Docker Image: a file consisting of all the dependencies needed to run our application
- Amazon Container Registry (ACR): Where we register the Docker Image we created
- Azure Web Apps (PAAS): Runs the Docker Container that we instantiated from the Docker Image

Example Workflow: Level III

Containerize and Deploy Machine Learning Pipeline on Google Kubernetes Engine




<https://towardsdatascience.com/deploy-machine-learning-model-on-google-kubernetes-engine-94daac85108b>

- Docker Image: a file consisting of all the dependencies needed to run our application
- Google Container Registry (ACR): Where we register the Docker Image we created
- Google Kubernetes Engine: Runs the Docker Container across a cluster which can autoscale

Example Workflow: Level III

Containerize and Deploy a Streamlit app on Google Kubernetes Engine



<https://towardsdatascience.com/deploy-machine-learning-app-built-using-streamlit-and-pycaret-on-google-kubernetes-engine-fd7e393d99cb>

Stats AI

- Note: Before we used Flask + HTML to deploy our GUI. Here instead Streamlit is used to do the same thing.
- Docker Image: a file consisting of all the dependencies needed to run our application
- Google Container Registry (ACR): Where we register the Docker Image we created
- Google Kubernetes Engine: Runs the Docker Container across a cluster which can autoscale

Example Workflow: Level III

Containerize and Deploy a Streamlit app on AWS Fargate



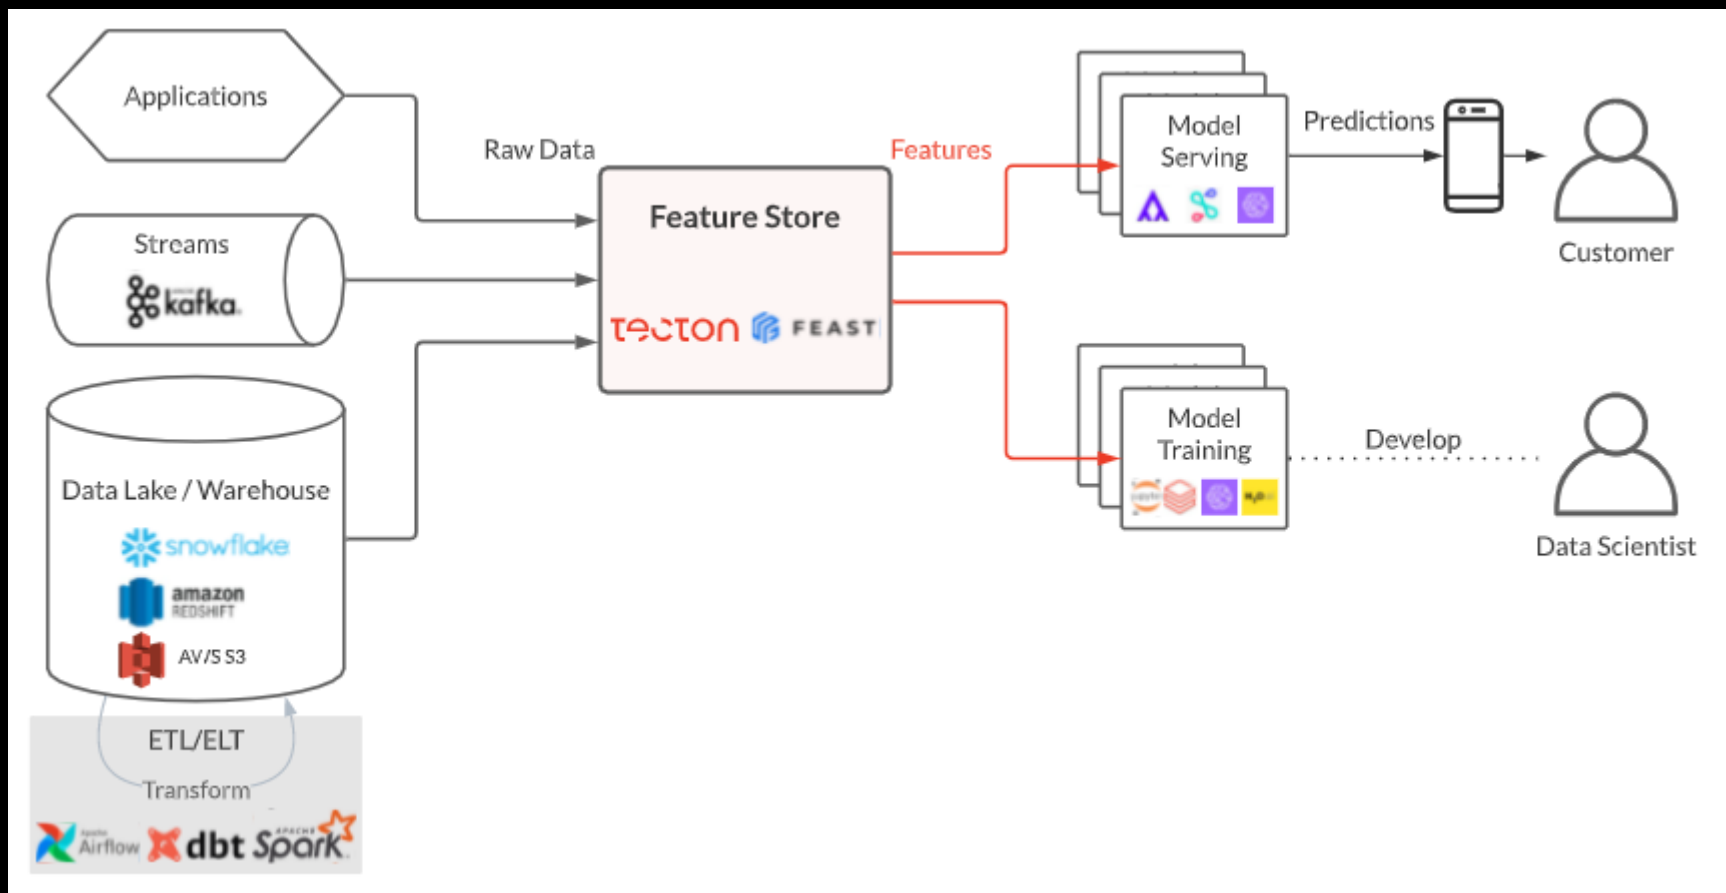
<https://towardsdatascience.com/deploy-pycaret-and-streamlit-app-using-aws-fargate-serverless-infrastructure-8b7d7c0584c2>

- AWS Fargate/Google Cloud Run are serverless PAAS (benefits of development w/o resource management)
- Docker Image: a file consisting of all the dependencies needed to run our application
- AWS Elastic Container Registry (EACR): Where we register the Docker Image we created
- AWS Fargate: Auto manages the resources on ECS or EKS to scale the Docker Container

The Rise of Feature Stores

Feature Stores

- Are basically databases for features
- By tightly coupling feature generation to data generating process, the Feature Store eliminates the possibility of finding features found in R&D that aren't possible to implement in production
- Eliminates Data Leakage:
 - Data is uploaded nightly.
 - The model developed in R&D finds that intraday data is best.
 - The problem is that that intraday data only exists after the fact.
 - Thus, only the prior night's data would be available for the production system
 - Thus, the feature store enables the data scientist developing the R&D system to “know” that intraday data would not be available at the time of inference
- Bottom line: The Feature stores value prop is its ability to serve as a single source of truth for development of both ML models and inference on fresh input values



Feature Stores

- Introduce economies of scale:
 - Features discovered for one model are often useful in other models
 - We can group features by the tasks they are used in
 - regression tasks, classification tasks, anomaly detection tasks
 - When a newly created feature is registered in a feature store, it becomes available for immediate reuse by every other model across the organization
 - This reduces duplication of data engineering efforts and allows new ML projects to bootstrap with a library of curated production-ready features

A feature store is an ML-specific data system that:

- Runs data pipelines that **transform** raw data into feature values
- **Stores** and manages the feature data itself, and
- **Serves** feature data consistently for training and inference purposes

